
CS6700 : Reinforcement Learning

Programming Assignment #1

Multi-armed Bandits

Deadline: 21st Feb 2018, 11:55 pm

- This is an individual assignment. Collaborations and discussions are strictly prohibited.
- You have to turn in the well-documented code along with a detailed report of the results of the experiments. Typeset your report in L^AT_EX.
- Be precise with your explanations. Unnecessary verbosity will be penalized.
- Check the Moodle discussion forums regularly for updates regarding the assignment.
- *Any kind* of plagiarism will be dealt with extremely seriously. Acknowledge any and every resource used.
- **Please start early.**

You are to conduct experiments on the 10-arm bandit testbed described in Section 2.3 of [the book](#). Please turn in the code for the testbed as well as the algorithms. Label the graphs clearly, with axes labels, parameter values, question numbers etc. Ensure that the code is adequately commented. Also turn in a short write-up of your observations from the experiments and answers to all the questions asked below.

Remember: The graphs are to run for 1000 plays, with each curve being the average of the performance on 2000 different bandit problems, generated as per the description in the book.

Points will be given according to the following criteria :

- Correct generation of the bandit problems.
- Correct coding of the learning algorithms.
- Correct code for gathering data to plot the graphs,
- Performance of the learning algorithms (correctness, optimality),
- Neatness of the graphs, appropriate labelling etc, and well-commented code.

Note: You can program in any language you want, but if you need any help later on, it will have to be in a language the TAs are comfortable with. So check with us before you start.

1. (10 marks) First implement the 10-arm bandit testbed. Try to reproduce the graphs given in Figure 2.2 on page 23 of the textbook. Specifically, implement the ϵ -greedy method on the 10-armed bandit testbed, for different values of ϵ .

Use graphing software of your choice, (for eg. Matlab or Gnuplot), to produce the graphs. Ensure that you have labelled the graphs correctly.

2. (8 marks) How does the softmax action selection method using the Gibbs distribution fare on the 10-armed testbed? Implement the method and run it at several (at least 3) temperatures to produce graphs similar to the previous one. Note that now you are required to sample from the softmax distribution, and not take the action with the maximum probability!
3. (7 marks) Implement the UCB1 algorithm. Compare the performance with that of epsilon greedy and softmax. What do you observe? Why do you think this is so?
4. (10 marks) Implement the Median Elimination Algorithm discussed in class.
 - Compare the performance with the others. What do you observe? Why do you think this is so?
 - Within MEA, find out if the finding the median is the rate-determining-step. Can you make it faster?
5. (5 marks) What happens as the number of arms grows? Run experiments on a 1000 arm bandit setup and compare.

Submission Guidelines

Submit a single tar/zip file containing the following files in the specified directory structure. Use the following naming convention: `rollno_PA1.tar.gz` or `rollno_PA1.zip`
A sample submission would look like this:

```
rollno_PA1
├── Code
│   ├── experiment_testbed_file(s)
│   ├── algorithm_file(s)
│   ├── plotting_file(s)
│   └── ...
├── report.pdf
└── README
```